

```
In [1]: 1 #importing all the libraries
2
3 import warnings
4 warnings.filterwarnings("ignore")
5 from sklearn.datasets import load_boston
6 from random import seed
7 from random import randrange
8 from csv import reader
9 from math import sqrt
10 from sklearn import preprocessing
11 import pandas as pd
12 import numpy as np
13 import matplotlib.pyplot as plt
14 from prettytable import PrettyTable
15 from sklearn.linear_model import SGDRegressor
16 from sklearn import preprocessing
17 from sklearn.metrics import mean_squared_error
```

Loading the data

```
In [2]: 1 data = load_boston()  
2 #data description  
3 print(data.DESCR)
```

```
Boston House Prices dataset  
=====
```

Notes

Data Set Characteristics:

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive

:Median Value (attribute 14) is usually the target

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<http://archive.ics.uci.edu/ml/datasets/Housing> (<http://archive.ics.uci.edu/ml/datasets/Housing>)

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

****References****

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.
- many more! (see <http://archive.ics.uci.edu/ml/datasets/Housing>) (<http://archive.ics.uci.edu/ml/datasets/Housing>)

```
In [3]: 1 data.feature_names#all the features in the dataset
```

```
Out[3]: array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',  
              'TAX', 'PTRATIO', 'B', 'LSTAT'],  
             dtype='<U7')
```

```
In [4]: 1 print('Shape of dataset is',data.data.shape)  
        2 print('shape of target variable is',data.target.shape)
```

```
Shape of dataset is (506, 13)  
shape of target variable is (506,)
```

```
In [5]: 1 boston_df = pd.DataFrame(data.data, columns = data.feature_names)
        2 boston_df.head()
```

Out[5]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [6]: 1 boston_df.describe()
```

Out[6]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRAT
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.593761	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.237154	18.4555
std	8.596783	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.537116	2.1649
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.000000	12.6000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.000000	17.4000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.000000	19.0500
75%	3.647423	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.000000	20.2000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000	22.0000

Splitting the data

```
In [7]: 1 X = data.data
        2 Y = data.target
```

```
In [8]: 1 #performing random splitting as data is not temporal in nature
2
3 from sklearn.model_selection import train_test_split
4 X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.2,random_state=42)
5 #X_train,X_val,Y_train,Y_val = train_test_split(X_train,Y_train,test_size = 0.1,random_state = 42)
6 print('after splitting :')
7 print('Size of training data',X_train.shape)
8 print('Size of test data',X_test.shape)
```

after splitting :

Size of training data (404, 13)

Size of test data (102, 13)

Standardizing the data

```
In [9]: 1 from sklearn.preprocessing import StandardScaler
2
3 std = StandardScaler().fit(X_train)
4 X_train = std.transform(X_train)
5 X_test = std.transform(X_test)
6
7 print(X_train.shape)
8 print(X_test.shape)
```

(404, 13)

(102, 13)

1.1 Implementing SKlearn SGDRegressor

```
In [10]: 1 from sklearn.linear_model import SGDRegressor
2
3 lr_model = SGDRegressor(loss = 'squared_loss', n_iter = 5000)
4 lr_model.fit(X_train, Y_train)
5
6 lr_model
```

```
Out[10]: SGDRegressor(alpha=0.0001, average=False, epsilon=0.1, eta0=0.01,
    fit_intercept=True, l1_ratio=0.15, learning_rate='invscaling',
    loss='squared_loss', n_iter=5000, penalty='l2', power_t=0.25,
    random_state=None, shuffle=True, verbose=0, warm_start=False)
```

```
In [11]: 1 #predicting on test data
2 Y_hat_skl = lr_model.predict(X_test)
3 print('Coefficients from this implementation are\n', lr_model.coef_)
4 print('=====\n')
5 print('optimal intercept is:', lr_model.intercept_)
6
7 mse_sklearn = mean_squared_error(Y_test, Y_hat_skl)
8 print('=====\n')
9 print('mean squared error from sklearn implementation is : ', mse_sklearn)
```

```
Coefficients from this implementation are
[-0.99701126  0.69511695  0.28151959  0.71766243 -2.0261877  3.14102981
-0.17216989 -3.08507052  2.24880257 -1.76001057 -2.0393385  1.13617232
-3.61922561]
```

```
=====
```

```
optimal intercept is: [ 22.79635486]
```

```
=====
```

```
mean squared error from sklearn implementation is : 24.3033972171
```

1.2 Implementing SGD Manually (from Scratch):

- While implementing SGD we will be doing two implementation trying with different learning rates.
 - learning rate changing with each epoch
 - learning rate constant with each epoch

1.2.1 Changing learning rate with each epoch

```
In [12]: 1 Weights = np.random.randn(X_train.shape[1],1)#initializing the weights
2 Coeff = np.random.randn(1,1)#initializing the coefficient values
3
4 n_epochs = 5000#number of times we want all of the data to be given as input
5 m = X_train.shape[0]#total number of data points in the training data
6 ep_loss = [] #list for loss at the end of each epoch
7 lr = 1 #learning rate
8
9 for ep in range(1,n_epochs+1):
10     loss = 0 #the mean squared error
11
12
13     for i in range(m):
14         #here we are looping over the total number of data points once,for completion of one epoch
15
16         b = np.random.randint(0,m)#the batch size ,as we want one random data point per iteration
17
18         #reshaping the one training sample and corresponding y
19         X = X_train[b,:].reshape(1,X_train.shape[1])
20         Y = Y_train[b].reshape(1,1)
21
22
23         Y_pred = np.dot(X,Weights) + Coeff #finding the predicted values on training data
24         loss = loss + (Y_pred - Y)**2 #computing the sum of mean squared errors in every iteration
25
26
27         Weights = Weights - (2/m)*(lr)*(X.T.dot(Y_pred - Y))#updating the weights
28         Coeff = Coeff - (2/m)*(lr)*(Y_pred - Y) #updating the bias term
29
30     lr = lr/2 # learning rate gets reduced to half with every epoch
31
32
33     ep_loss.append(loss[0][0]/m)
34
```

```

In [13]: 1 weights_vlr = Weights#value of weight vector at the end of 5000th epoch
2         Coeff_vlr = Coeff
3         print('Weights for learning rate changing with every epoch are:',Weights)
4         print('Constant term for learning rate changing with every epoch is:',Coeff)
5
6         %matplotlib inline
7
8         #plotting the epochs vs loss
9         plt.figure(figsize = (12,6))
10        plt.plot([i for i in range(0,n_epochs)],ep_loss)
11        plt.xlabel('Epochs',fontsize = 15)
12        plt.ylabel('loss',fontsize = 15)
13        plt.title('epochs vs loss with changing learning rate',fontsize = 15)
14
15        Y_hat = ((Weights.T).dot(X_test.T) + Coeff ).reshape(102,1)#computing the predicted values on test data
16        Y_test = np.array(Y_test)
17        mse_sgd = mean_squared_error(Y_test,Y_hat)#computing the mean squared error
18        print('mean_squared_error after implementing sgd manually with changing leaarning rate is :',mse_sgd)
19


```

Weights for learning rate changing with every epoch are: $\begin{bmatrix} -0.78454833 \\ 0.3468268 \\ -0.03658473 \\ 1.04243126 \\ -0.97777929 \\ 3.50226501 \\ -0.21894241 \\ -1.78534528 \\ 1.04669725 \\ -1.24923678 \\ -1.58278309 \\ 1.10452903 \\ -3.10811876 \end{bmatrix}$

$\begin{bmatrix} 0.3468268 \\ -0.03658473 \\ 1.04243126 \\ -0.97777929 \\ 3.50226501 \\ -0.21894241 \\ -1.78534528 \\ 1.04669725 \\ -1.24923678 \\ -1.58278309 \\ 1.10452903 \\ -3.10811876 \end{bmatrix}$

Constant term for learning rate changing with every epoch is: $\begin{bmatrix} 22.42254032 \end{bmatrix}$

mean_squared_error after implementing sgd manually with changing leaarning rate is : 26.6180348135



1.2.2 Constant learning rate with each epoch

```

In [14]: 1 Weights = np.random.randn(X_train.shape[1],1)#initializing the weights
2 Coeff = np.random.randn(1,1)#initializing the coefficient values
3
4 n_epochs = 5000#number of times we want all of the data to be given as input
5 m = X_train.shape[0]#total number of data points in the training data
6 ep_loss = [] #list for loss at the end of each epoch
7 lr = 1 #learning rate
8
9 for ep in range(1,n_epochs+1):
10     loss = 0 #the mean squared error
11
12
13     for i in range(m):
14         #here we are looping over the total number of data points once,for completion of one epoch
15
16         b = np.random.randint(0,m)#the batch size ,as we want one random data point per iteration
17
18         #reshaping the one training sample and corresponding y
19         X = X_train[b,:].reshape(1,X_train.shape[1])
20         Y = Y_train[b].reshape(1,1)
21
22
23         Y_pred = np.dot(X,Weights) + Coeff #finding the predicted values on training data
24         loss = loss + (Y_pred - Y)**2 #computing the sum of mean squared errors in every iteration
25
26
27         Weights = Weights - (2/m)*(lr)*(X.T.dot(Y_pred - Y))#updating the weights
28         Coeff = Coeff - (2/m)*(lr)*(Y_pred - Y) #updating the bias term
29
30     #Lr = Lr/2
31
32
33     ep_loss.append(loss[0][0]/m)
34

```

```

In [16]: 1 weights_c = Weights#value of weight vector at the end of 5000th epoch
2         Coeff_c = Coeff
3         print('Weights for learning rate changing with every epoch are:',Weights)
4         print('Constant term for learning rate changing with every epoch is:',Coeff)
5
6         %matplotlib inline
7
8         #plotting the epochs vs loss
9         plt.figure(figsize = (12,6))
10        plt.plot([i for i in range(0,n_epochs)],ep_loss)
11        plt.xlabel('Epochs',fontsize = 15)
12        plt.ylabel('loss',fontsize = 15)
13        plt.title('epochs vs loss with constant learning rate',fontsize = 15)
14
15        Y_hat_c = ((Weights.T).dot(X_test.T) + Coeff ).reshape(102,1)#computing the predicted values on test data
16        Y_test = np.array(Y_test)
17        mse_sgd_c = mean_squared_error(Y_test,Y_hat_c)#computing the mean squared error
18        print('mean_squared_error after implementing sgd manually with constant learning rate is :',mse_sgd_c)
19

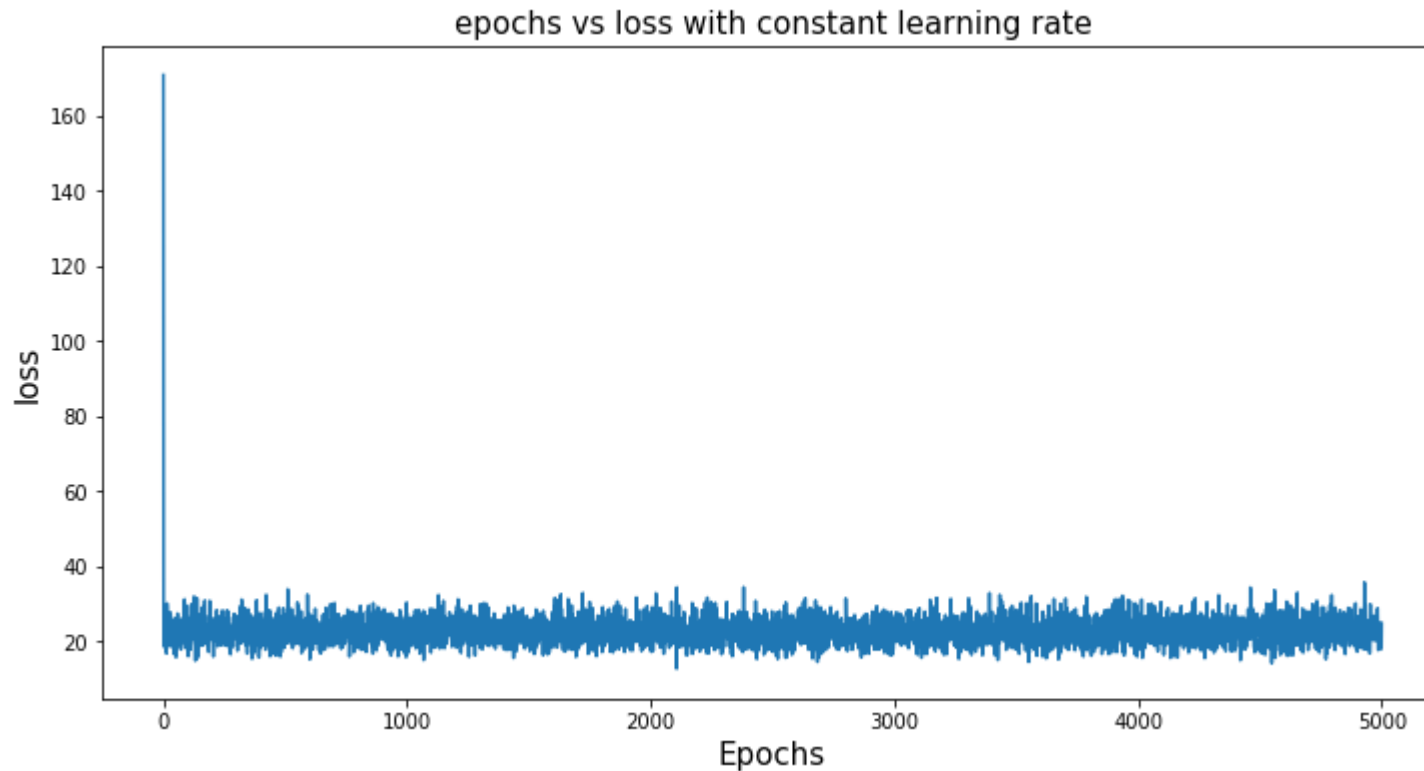
```

Weights for learning rate changing with every epoch are: $\begin{bmatrix} -0.9872763 \end{bmatrix}$

$\begin{bmatrix} 0.85009917 \\ 0.06636749 \\ 0.87204787 \\ -2.11683567 \\ 3.15769378 \\ -0.52586822 \\ -3.37107063 \\ 2.55775089 \\ -1.71615317 \\ -1.83617658 \\ 0.72633639 \\ -3.26997752 \end{bmatrix}$

Constant term for learning rate changing with every epoch is: $\begin{bmatrix} 23.21874236 \end{bmatrix}$

mean_squared_error after implementing sgd manually with constant learning rate is : 22.8566063721



We see that the loss converges very fast between 0 to 10 number of epochs and then it oscillates between a certain minimum and maximum ,as the function is trying to find a global minima for convergence.

Comparison between Sklearn and Manual Implementation


In [17]:

```
1 import seaborn as sns
2
3
4 sns.set_style('darkgrid')
5 plt.figure(figsize = (10,6))
6 plt.subplot(211)
7 plt.scatter(Y_test,Y_hat_skl)
8 plt.plot(np.linspace(0,50,10),np.linspace(0,50,10),'r')
9 plt.xlabel("Prices:  $Y_i$ ",fontsize = 15)
10 plt.ylabel("Predicted prices:  $\hat{Y}_i$ ",fontsize = 15)
11 plt.title("Prices vs Predicted prices with Sklearn implementation:  $Y_i$  vs  $\hat{Y}_i$ ",fontsize = 15)
12 #plt.show()
13
14
15 plt.figure(figsize = (10,6))
16 plt.subplot(212)
17 plt.scatter(Y_test,Y_hat)
18 plt.plot(np.linspace(0,50,10),np.linspace(0,50,10),'r')
19 plt.xlabel("Prices:  $Y_i$ ",fontsize = 15)
20 plt.ylabel("Predicted prices:  $\hat{Y}_i$ ",fontsize = 15)
21 plt.title("Prices vs Predicted prices with SGD implementation(changing learning rate):  $Y_i$  vs  $\hat{Y}_i$ ")
22 plt.show()
23
24
25 plt.figure(figsize = (10,6))
26 plt.subplot(212)
27 plt.scatter(Y_test,Y_hat_c)
28 plt.plot(np.linspace(0,50,10),np.linspace(0,50,10),'r')
29 plt.xlabel("Prices:  $Y_i$ ",fontsize = 15)
30 plt.ylabel("Predicted prices:  $\hat{Y}_i$ ",fontsize = 15)
31 plt.title("Prices vs Predicted prices with SGD implementation(constant learning rate):  $Y_i$  vs  $\hat{Y}_i$ ")
32 plt.show()
33
34
```





```
In [54]: 1 """WE will plot the kernel density estimation plots of error and predicted values"""
2
3 #differenc between actual and predicted values for all three cases
4 delta_y1 = Y_test - Y_hat_skl #sklearn
5 delta_y2 = np.asmatrix(Y_test) - Y_hat #changing learning rate
6 delta_y3 = np.asmatrix(Y_test) - Y_hat_c #constant learning rate
7
8 sns.set_style('darkgrid')
9 plt.figure(figsize = (12,5))
10 sns.kdeplot(np.array(delta_y1), bw=0.5,label = 'SKlearn')
11 sns.kdeplot(np.asarray(delta_y2)[0], bw=0.5,label = 'Manual with changing learning rate')
12 sns.kdeplot(np.asarray(delta_y3)[0], bw=0.5,label = 'Manual with Constant learning rate')
13 plt.title('Distribution of error',fontsize = 15)
14 plt.xlabel('Error',fontsize = 15)
15 plt.ylabel('Density',fontsize = 15)
16 plt.legend(loc = 'best',fontsize = 15)
17 plt.show()
18
19
20
21 #distribution density of predicted values
22
23 plt.figure(figsize = (12,5))
24 sns.kdeplot(np.asarray(Y_hat_skl), bw=0.5,label = 'SKlearn')
25 sns.kdeplot(np.asarray(Y_hat).T[0], bw=0.5,label = 'Manual with changing learning rate')
26 sns.kdeplot(np.asarray(Y_hat_c).T[0], bw=0.5,label = 'Manual with Constant learning rate')
27 plt.title('Distribution of Predicted values',fontsize = 15)
28 plt.xlabel('Predicted values',fontsize = 15)
29 plt.ylabel('Density',fontsize = 15)
30 plt.legend(loc = 'best',fontsize = 15)
31 plt.show()
32
33
34
35
```

```

1 - The distribution of error for Manual implemetation with changing learning rate has more variance than
  one with constant learning rate which can be understood with non constant lr.
2
3 -

```

```

In [76]: 1 # Comparing the final weights for each of the three impenetations
          2 weights_dict = {'SKlearn_implementation':lr_model.coef_, 'Manual with Changing lr':np.array(weights_vlr).T[0]
          3                  'Manual with constant LR':np.array(weights_c).T[0]}
          4 Weights_df = pd.DataFrame(weights_dict,index = None)
          5 Weights_df

```

Out[76]:

	SKlearn_implementation	Manual with Changing lr	Manual with constant LR
0	-0.997011	-0.784548	-0.987276
1	0.695117	0.346827	0.850099
2	0.281520	-0.036585	0.066367
3	0.717662	1.042431	0.872048
4	-2.026188	-0.977779	-2.116836
5	3.141030	3.502265	3.157694
6	-0.172170	-0.218942	-0.525868
7	-3.085071	-1.785345	-3.371071
8	2.248803	1.046697	2.557751
9	-1.760011	-1.249237	-1.716153
10	-2.039338	-1.582783	-1.836177
11	1.136172	1.104529	0.726336
12	-3.619226	-3.108119	-3.269978

```

In [78]: 1 print('MSE with SKlearn: ',mse_sklearn)
          2 print('MSE with Manual Implementation with Changing lr:',mse_sgd)
          3 print('MSE with Manual Implemetation with Constant lr:',mse_sgd_c)

```

MSE with SKlearn: 24.3033972171

MSE with Manual Implementation with Changing lr: 26.6180348135

MSE with Manual Implemetation with Constant lr: 22.8566063721

Conclusion:

- We get the best Mean Squared Error value with manual implementation of SGD keeping the learning rate constant for every epoch.
- The Sklearn implementation fared well than manual sgd where we changed the learning rate with every epoch. This can be attributed to the regularization terms which were used for avoiding overfitting in the data.
- Though None of the above model is good enough for prediction of regression values